

# Estimate generation v2

Use advanced models for generation estimation in the Global Power Plant Database. Primary model is a two-hidden-layer neural network.

In [1]:

```
# import what we'll need and set parameters

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from keras.models import Sequential
from keras.layers import Flatten, Dense, Lambda
from keras.layers import Conv2D, Dropout, Activation, MaxPooling2D
from keras.callbacks import EarlyStopping, ModelCheckpoint
from keras.utils.vis_utils import model_to_dot
from IPython.display import SVG
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
import pydot

GPPD_FILENAME = '../..//output_database/global_power_plant_database.csv'
WEIGHTS_FILE = "model/estimate_generation.h5"
VALIDATION_FRACTION = 0.2
```

Using TensorFlow backend.

In [2]:

```
# read in database
df = pd.read_csv(GPPD_FILENAME)
df.head()
```

Out[ 2 ] :

	country	country_long	name	gppd_idnr	capacity_mw	latitude	longitu
0	AFG	Afghanistan	Kajaki Hydroelectric Power Plant Afghanistan	GEODB0040538	33.00	32.3220	65.1190
1	AFG	Afghanistan	Mahipar Hydroelectric Power Plant Afghanistan	GEODB0040541	66.00	34.5560	69.4787
2	AFG	Afghanistan	Naghlu Dam Hydroelectric Power Plant Afghanistan	GEODB0040534	100.00	34.6410	69.7170
3	AFG	Afghanistan	Nangarhar (Darunta) Hydroelectric Power Plant ...	GEODB0040536	11.55	34.4847	70.3633
4	AFG	Afghanistan	Northwest Kabul Power Plant Afghanistan	GEODB0040540	42.00	34.5638	69.1134

5 rows × 22 columns

In [3]:

```
# show count for number of valid entries in each column  
df.count()
```

Out[3]:

country	25657
country_long	25657
name	25637
gppd_idnr	25657
capacity_mw	25657
latitude	25657
longitude	25657
fuel1	25657
fuel2	1670
fuel3	295
fuel4	107
commissioning_year	13933
owner	17157
source	25657
url	25657
geolocation_source	25657
year_of_capacity_data	16065
generation_gwh_2013	56
generation_gwh_2014	55
generation_gwh_2015	536
generation_gwh_2016	8657
estimated_generation_gwh	24941

dtype: int64

In [4]:

```
# prepare data for training

# convert string-type columns to categories (assume no NaNs in these columns)
factorized_countries,country_key = df['country'].factorize()
df['country'] = factorized_countries
factorized_fuel1,fuel1_key = df['fuel1'].factorize()
df['fuel1'] = factorized_fuel1

# convert numerical columns to np array to use as predictor variables and remove NaNs
X_columns = ['country','capacity_mw','latitude','longitude','commissioning_year',
            'fuel1']
df_No_NaN = df[X_columns + ['generation_gwh_2016']].dropna(how='any')
X_data = df_No_NaN[X_columns].as_matrix()

# calculate capacity factor to use as predicted variable
df_No_NaN['capacity_factor'] = df_No_NaN.apply(lambda row:row['generation_gwh_2016']/
(24.0*365.0*0.001*row['capacity_mw']),axis=1)
y_column = ['capacity_factor']
y_data = df_No_NaN[y_column].as_matrix()

# show results
print(X_data)
print(y_data)
print(len(X_data))
print(len(y_data))
```

```

[[ 8.00000000e+00  2.89000000e+02  4.72078000e+01  1.10057000e+0
1
    1.98100000e+03  0.00000000e+00]
 [ 8.00000000e+00  5.00000000e+02  4.72696000e+01  1.09678000e+0
1
    1.98100000e+03  0.00000000e+00]
 [ 4.30000000e+01  2.25000000e+03  3.02483000e+01  3.09471000e+0
1
    2.01400000e+03  1.00000000e+00]
 ...,
 [ 1.57000000e+02  2.80000000e+01  1.43611000e+01  1.08720300e+0
2
    2.01400000e+03  0.00000000e+00]
 [ 1.57000000e+02  1.95000000e+01  1.21526000e+01  1.08378700e+0
2
    2.01000000e+03  0.00000000e+00]
 [ 1.57000000e+02  3.00000000e+01  1.58600000e+01  1.07653800e+0
2
    2.00900000e+03  0.00000000e+00]]
[[ 0.04692255]
 [ 0.02934475]
 [ 0.00674784]
 ...,
 [ 0.41992825]
 [ 0.46247512]
 [ 0.46689498]]
8569
8569

```

In [5]:

```

# calculate scaling values for input data

mean_vals = np.mean(X_data,axis=0)
range_vals = np.max(X_data,axis=0) - np.min(X_data,axis=0)

```

In [9]:

```
# set up neural network

INPUT_SHAPE = X_data[0].shape
print(u"Input shape is: {0}".format(INPUT_SHAPE))
DROPOUT_RATE = 0.15
DENSE_LAYER_SIZE = 128

def myNet():
    model = Sequential()
    model.add(Lambda(lambda x: x - mean_vals,input_shape = INPUT_SHAPE))    # placeholder for normalization
    model.add(Dense(DENSE_LAYER_SIZE,activation='relu'))
    model.add(Dropout(DROPOUT_RATE))
    model.add(Dense(DENSE_LAYER_SIZE,activation='relu'))
    model.add(Dropout(DROPOUT_RATE))
    model.add(Dense(DENSE_LAYER_SIZE,activation='relu'))
    model.add(Dense(1))
    return model

model = myNet()
model.compile(loss='mean_squared_error',optimizer='adam',metrics=[ 'mean_absolute_error'])
print("Model contains {0} parameters.".format(model.count_params()))
print(model.summary())
```

Input shape is: (6,)  
Model contains 34049 parameters.

Layer (type)	Output Shape	Param #
=====		
lambda_3 (Lambda)	(None, 6)	0
<hr/>		
dense_9 (Dense)	(None, 128)	896
<hr/>		
dropout_5 (Dropout)	(None, 128)	0
<hr/>		
dense_10 (Dense)	(None, 128)	16512
<hr/>		
dropout_6 (Dropout)	(None, 128)	0
<hr/>		
dense_11 (Dense)	(None, 128)	16512
<hr/>		
dense_12 (Dense)	(None, 1)	129
=====		
Total params: 34,049		
Trainable params: 34,049		
Non-trainable params: 0		
<hr/>		
None		

In [10]:

```
# fit model

BATCH_SIZE = 64
NUM_EPOCHS = 128
early_stop = EarlyStopping(monitor='val_loss',min_delta=0.001,patience=32)
check_point = ModelCheckpoint(WEIGHTS_FILE,monitor='val_loss',save_best_only=True,mode='max')
history_object = model.fit(x=X_data, y=y_data,
                           batch_size = BATCH_SIZE,
                           epochs = NUM_EPOCHS,
                           verbose = 1,
                           callbacks = [early_stop,check_point],
                           validation_split = VALIDATION_FRACTION)

# reload model with best weights from training
model = myNet()
model.load_weights(WEIGHTS_FILE)
model.compile(loss='mean_squared_error',optimizer='adam',metrics=['mean_absolute_error'])
print("Finished training; model reloaded with optimum weights.")
model.save(WEIGHTS_FILE)
```

Train on 6855 samples, validate on 1714 samples

Epoch 1/128

6855/6855 [=====] - 1s 148us/step - loss: 14.4369 - mean\_absolute\_error: 3.7769 - val\_loss: 3.2402 - val\_mean\_absolute\_error: 0.8218

Epoch 2/128

6855/6855 [=====] - 1s 74us/step - loss: 20.1228 - mean\_absolute\_error: 1.9398 - val\_loss: 1.3747 - val\_mean\_absolute\_error: 0.6320

Epoch 3/128

6855/6855 [=====] - 1s 78us/step - loss: 7.7407 - mean\_absolute\_error: 1.2016 - val\_loss: 0.6293 - val\_mean\_absolute\_error: 0.5408

Epoch 4/128

6855/6855 [=====] - 1s 77us/step - loss: 7.0240 - mean\_absolute\_error: 1.0554 - val\_loss: 0.2429 - val\_mean\_absolute\_error: 0.3327

Epoch 5/128

6855/6855 [=====] - 1s 78us/step - loss: 3.2501 - mean\_absolute\_error: 0.8389 - val\_loss: 0.5816 - val\_mean\_absolute\_error: 0.4220

Epoch 6/128

6855/6855 [=====] - 1s 73us/step - loss: 2.2883 - mean\_absolute\_error: 0.7260 - val\_loss: 0.1865 - val\_mean\_absolute\_error: 0.3000

Epoch 7/128

6855/6855 [=====] - 1s 73us/step - loss: 2.1698 - mean\_absolute\_error: 0.6740 - val\_loss: 0.4389 - val\_mean\_absolute\_error: 0.3618

Epoch 8/128  
6855/6855 [=====] - 1s 74us/step - loss: 1.7815 - mean\_absolute\_error: 0.6200 - val\_loss: 0.3584 - val\_mean\_absolute\_error: 0.3555

Epoch 9/128  
6855/6855 [=====] - 0s 67us/step - loss: 1.6054 - mean\_absolute\_error: 0.5657 - val\_loss: 0.3719 - val\_mean\_absolute\_error: 0.3657

Epoch 10/128  
6855/6855 [=====] - 1s 77us/step - loss: 1.5256 - mean\_absolute\_error: 0.5457 - val\_loss: 0.1796 - val\_mean\_absolute\_error: 0.2860

Epoch 11/128  
6855/6855 [=====] - 1s 81us/step - loss: 0.9885 - mean\_absolute\_error: 0.4782 - val\_loss: 0.3013 - val\_mean\_absolute\_error: 0.3249

Epoch 12/128  
6855/6855 [=====] - 0s 73us/step - loss: 0.9233 - mean\_absolute\_error: 0.4612 - val\_loss: 0.1798 - val\_mean\_absolute\_error: 0.2878

Epoch 13/128  
6855/6855 [=====] - 0s 66us/step - loss: 0.8753 - mean\_absolute\_error: 0.4530 - val\_loss: 0.1124 - val\_mean\_absolute\_error: 0.2656

Epoch 14/128  
6855/6855 [=====] - 1s 77us/step - loss: 0.5310 - mean\_absolute\_error: 0.3890 - val\_loss: 0.0920 - val\_mean\_absolute\_error: 0.2378

Epoch 15/128  
6855/6855 [=====] - 1s 83us/step - loss: 0.5966 - mean\_absolute\_error: 0.3825 - val\_loss: 0.1926 - val\_mean\_absolute\_error: 0.2805

Epoch 16/128  
6855/6855 [=====] - 1s 80us/step - loss: 0.4950 - mean\_absolute\_error: 0.3657 - val\_loss: 0.0735 - val\_mean\_absolute\_error: 0.2255

Epoch 17/128  
6855/6855 [=====] - 1s 75us/step - loss: 0.4385 - mean\_absolute\_error: 0.3537 - val\_loss: 0.1095 - val\_mean\_absolute\_error: 0.2615

Epoch 18/128  
6855/6855 [=====] - 1s 88us/step - loss: 0.3948 - mean\_absolute\_error: 0.3332 - val\_loss: 0.1518 - val\_mean\_absolute\_error: 0.2626

Epoch 19/128  
6855/6855 [=====] - 1s 95us/step - loss: 0.3311 - mean\_absolute\_error: 0.3197 - val\_loss: 0.1052 - val\_mean\_absolute\_error: 0.2358

Epoch 20/128  
6855/6855 [=====] - 1s 81us/step - loss: 0.3156 - mean\_absolute\_error: 0.3168 - val\_loss: 0.0680 - val\_mean\_absolute\_error: 0.2082

Epoch 21/128



6855/6855 [=====] - 1s 78us/step - loss: 0.3840 - mean\_absolute\_error: 0.3080 - val\_loss: 0.1658 - val\_mean\_absolute\_error: 0.2511  
Epoch 22/128  
6855/6855 [=====] - 1s 81us/step - loss: 0.3782 - mean\_absolute\_error: 0.3240 - val\_loss: 0.0949 - val\_mean\_absolute\_error: 0.2318  
Epoch 23/128  
6855/6855 [=====] - 1s 78us/step - loss: 0.2716 - mean\_absolute\_error: 0.2948 - val\_loss: 0.0829 - val\_mean\_absolute\_error: 0.2132  
Epoch 24/128  
6855/6855 [=====] - 1s 76us/step - loss: 0.2109 - mean\_absolute\_error: 0.2753 - val\_loss: 0.0642 - val\_mean\_absolute\_error: 0.2014  
Epoch 25/128  
6855/6855 [=====] - 0s 73us/step - loss: 0.1809 - mean\_absolute\_error: 0.2659 - val\_loss: 0.0638 - val\_mean\_absolute\_error: 0.1976  
Epoch 26/128  
6855/6855 [=====] - 1s 75us/step - loss: 0.1806 - mean\_absolute\_error: 0.2641 - val\_loss: 0.0623 - val\_mean\_absolute\_error: 0.1903  
Epoch 27/128  
6855/6855 [=====] - 1s 75us/step - loss: 0.2145 - mean\_absolute\_error: 0.2650 - val\_loss: 0.0641 - val\_mean\_absolute\_error: 0.1999  
Epoch 28/128  
6855/6855 [=====] - 1s 74us/step - loss: 0.1701 - mean\_absolute\_error: 0.2625 - val\_loss: 0.0633 - val\_mean\_absolute\_error: 0.2014  
Epoch 29/128  
6855/6855 [=====] - 1s 77us/step - loss: 0.1641 - mean\_absolute\_error: 0.2528 - val\_loss: 0.0656 - val\_mean\_absolute\_error: 0.2034  
Epoch 30/128  
6855/6855 [=====] - 1s 76us/step - loss: 0.1418 - mean\_absolute\_error: 0.2491 - val\_loss: 0.0625 - val\_mean\_absolute\_error: 0.1972  
Epoch 31/128  
6855/6855 [=====] - 1s 80us/step - loss: 0.1218 - mean\_absolute\_error: 0.2374 - val\_loss: 0.0626 - val\_mean\_absolute\_error: 0.1930  
Epoch 32/128  
6855/6855 [=====] - 1s 80us/step - loss: 0.1265 - mean\_absolute\_error: 0.2388 - val\_loss: 0.0662 - val\_mean\_absolute\_error: 0.2055  
Epoch 33/128  
6855/6855 [=====] - 1s 77us/step - loss: 0.1345 - mean\_absolute\_error: 0.2413 - val\_loss: 0.0789 - val\_mean\_absolute\_error: 0.2135  
Epoch 34/128  
6855/6855 [=====] - 1s 79us/step - loss: 0.


1251 - mean\_absolute\_error: 0.2352 - val\_loss: 0.0637 - val\_mean\_absolute\_error: 0.2006  
Epoch 35/128  
6855/6855 [=====] - 1s 80us/step - loss: 0.1502 - mean\_absolute\_error: 0.2431 - val\_loss: 0.0678 - val\_mean\_absolute\_error: 0.1999  
Epoch 36/128  
6855/6855 [=====] - 1s 80us/step - loss: 0.1117 - mean\_absolute\_error: 0.2302 - val\_loss: 0.0619 - val\_mean\_absolute\_error: 0.1942  
Epoch 37/128  
6855/6855 [=====] - 1s 75us/step - loss: 0.1240 - mean\_absolute\_error: 0.2275 - val\_loss: 0.0650 - val\_mean\_absolute\_error: 0.1971  
Epoch 38/128  
6855/6855 [=====] - 1s 76us/step - loss: 0.0970 - mean\_absolute\_error: 0.2241 - val\_loss: 0.0625 - val\_mean\_absolute\_error: 0.1968  
Epoch 39/128  
6855/6855 [=====] - 1s 73us/step - loss: 0.1294 - mean\_absolute\_error: 0.2278 - val\_loss: 0.0764 - val\_mean\_absolute\_error: 0.2117  
Epoch 40/128  
6855/6855 [=====] - 1s 74us/step - loss: 0.1140 - mean\_absolute\_error: 0.2253 - val\_loss: 0.0648 - val\_mean\_absolute\_error: 0.2023  
Epoch 41/128  
6855/6855 [=====] - 0s 72us/step - loss: 0.0967 - mean\_absolute\_error: 0.2202 - val\_loss: 0.0641 - val\_mean\_absolute\_error: 0.2019  
Epoch 42/128  
6855/6855 [=====] - 1s 76us/step - loss: 0.0839 - mean\_absolute\_error: 0.2145 - val\_loss: 0.0622 - val\_mean\_absolute\_error: 0.1926  
Epoch 43/128  
6855/6855 [=====] - 1s 76us/step - loss: 0.0851 - mean\_absolute\_error: 0.2135 - val\_loss: 0.0633 - val\_mean\_absolute\_error: 0.1953  
Epoch 44/128  
6855/6855 [=====] - 1s 75us/step - loss: 0.0853 - mean\_absolute\_error: 0.2124 - val\_loss: 0.0633 - val\_mean\_absolute\_error: 0.2004  
Epoch 45/128  
6855/6855 [=====] - 0s 73us/step - loss: 0.0842 - mean\_absolute\_error: 0.2128 - val\_loss: 0.0635 - val\_mean\_absolute\_error: 0.2003  
Epoch 46/128  
6855/6855 [=====] - 0s 73us/step - loss: 0.0766 - mean\_absolute\_error: 0.2088 - val\_loss: 0.0775 - val\_mean\_absolute\_error: 0.2097  
Epoch 47/128  
6855/6855 [=====] - 1s 75us/step - loss: 0.0779 - mean\_absolute\_error: 0.2102 - val\_loss: 0.0647 - val\_mean\_absolute\_error: 0.2097

olute\_error: 0.1997  
Epoch 48/128  
6855/6855 [=====] - 1s 75us/step - loss: 0.0730 - mean\_absolute\_error: 0.2072 - val\_loss: 0.0624 - val\_mean\_absolute\_error: 0.1963  
Epoch 49/128  
6855/6855 [=====] - 1s 87us/step - loss: 0.0713 - mean\_absolute\_error: 0.2047 - val\_loss: 0.0624 - val\_mean\_absolute\_error: 0.1969  
Epoch 50/128  
6855/6855 [=====] - 0s 71us/step - loss: 0.0742 - mean\_absolute\_error: 0.2051 - val\_loss: 0.0625 - val\_mean\_absolute\_error: 0.1971  
Epoch 51/128  
6855/6855 [=====] - 0s 68us/step - loss: 0.0696 - mean\_absolute\_error: 0.2028 - val\_loss: 0.0647 - val\_mean\_absolute\_error: 0.1991  
Epoch 52/128  
6855/6855 [=====] - 1s 75us/step - loss: 0.0667 - mean\_absolute\_error: 0.2010 - val\_loss: 0.0652 - val\_mean\_absolute\_error: 0.1995  
Epoch 53/128  
6855/6855 [=====] - 1s 82us/step - loss: 0.0688 - mean\_absolute\_error: 0.2014 - val\_loss: 0.0639 - val\_mean\_absolute\_error: 0.1989  
Epoch 54/128  
6855/6855 [=====] - 1s 79us/step - loss: 0.0660 - mean\_absolute\_error: 0.2012 - val\_loss: 0.0616 - val\_mean\_absolute\_error: 0.1963  
Epoch 55/128  
6855/6855 [=====] - 1s 79us/step - loss: 0.0699 - mean\_absolute\_error: 0.2017 - val\_loss: 0.0622 - val\_mean\_absolute\_error: 0.1960  
Epoch 56/128  
6855/6855 [=====] - 1s 80us/step - loss: 0.0769 - mean\_absolute\_error: 0.2046 - val\_loss: 0.0628 - val\_mean\_absolute\_error: 0.1966  
Epoch 57/128  
6855/6855 [=====] - ETA: 0s - loss: 0.0666 - mean\_absolute\_error: 0.199 - 1s 85us/step - loss: 0.0665 - mean\_absolute\_error: 0.1997 - val\_loss: 0.0628 - val\_mean\_absolute\_error: 0.1965  
Epoch 58/128  
6855/6855 [=====] - 1s 81us/step - loss: 0.0669 - mean\_absolute\_error: 0.1984 - val\_loss: 0.0639 - val\_mean\_absolute\_error: 0.1984  
Finished training; model reloaded with optimum weights.

In [11]:

```
# plot training loss history
```

```
plt.plot(history_object.history['mean_absolute_error'])  
plt.plot(history_object.history['val_loss'])  
plt.ylabel('loss')  
plt.xlabel('epoch')  
plt.legend(['training loss (mse)', 'validation loss'], loc='upper right')  
plt.ylim([0, 0.5])  
plt.show()
```



In [12]:

```
# visualize model
```

```
SVG(model_to_dot(model).create(prog='dot', format='svg'))
```

Out[12]:

